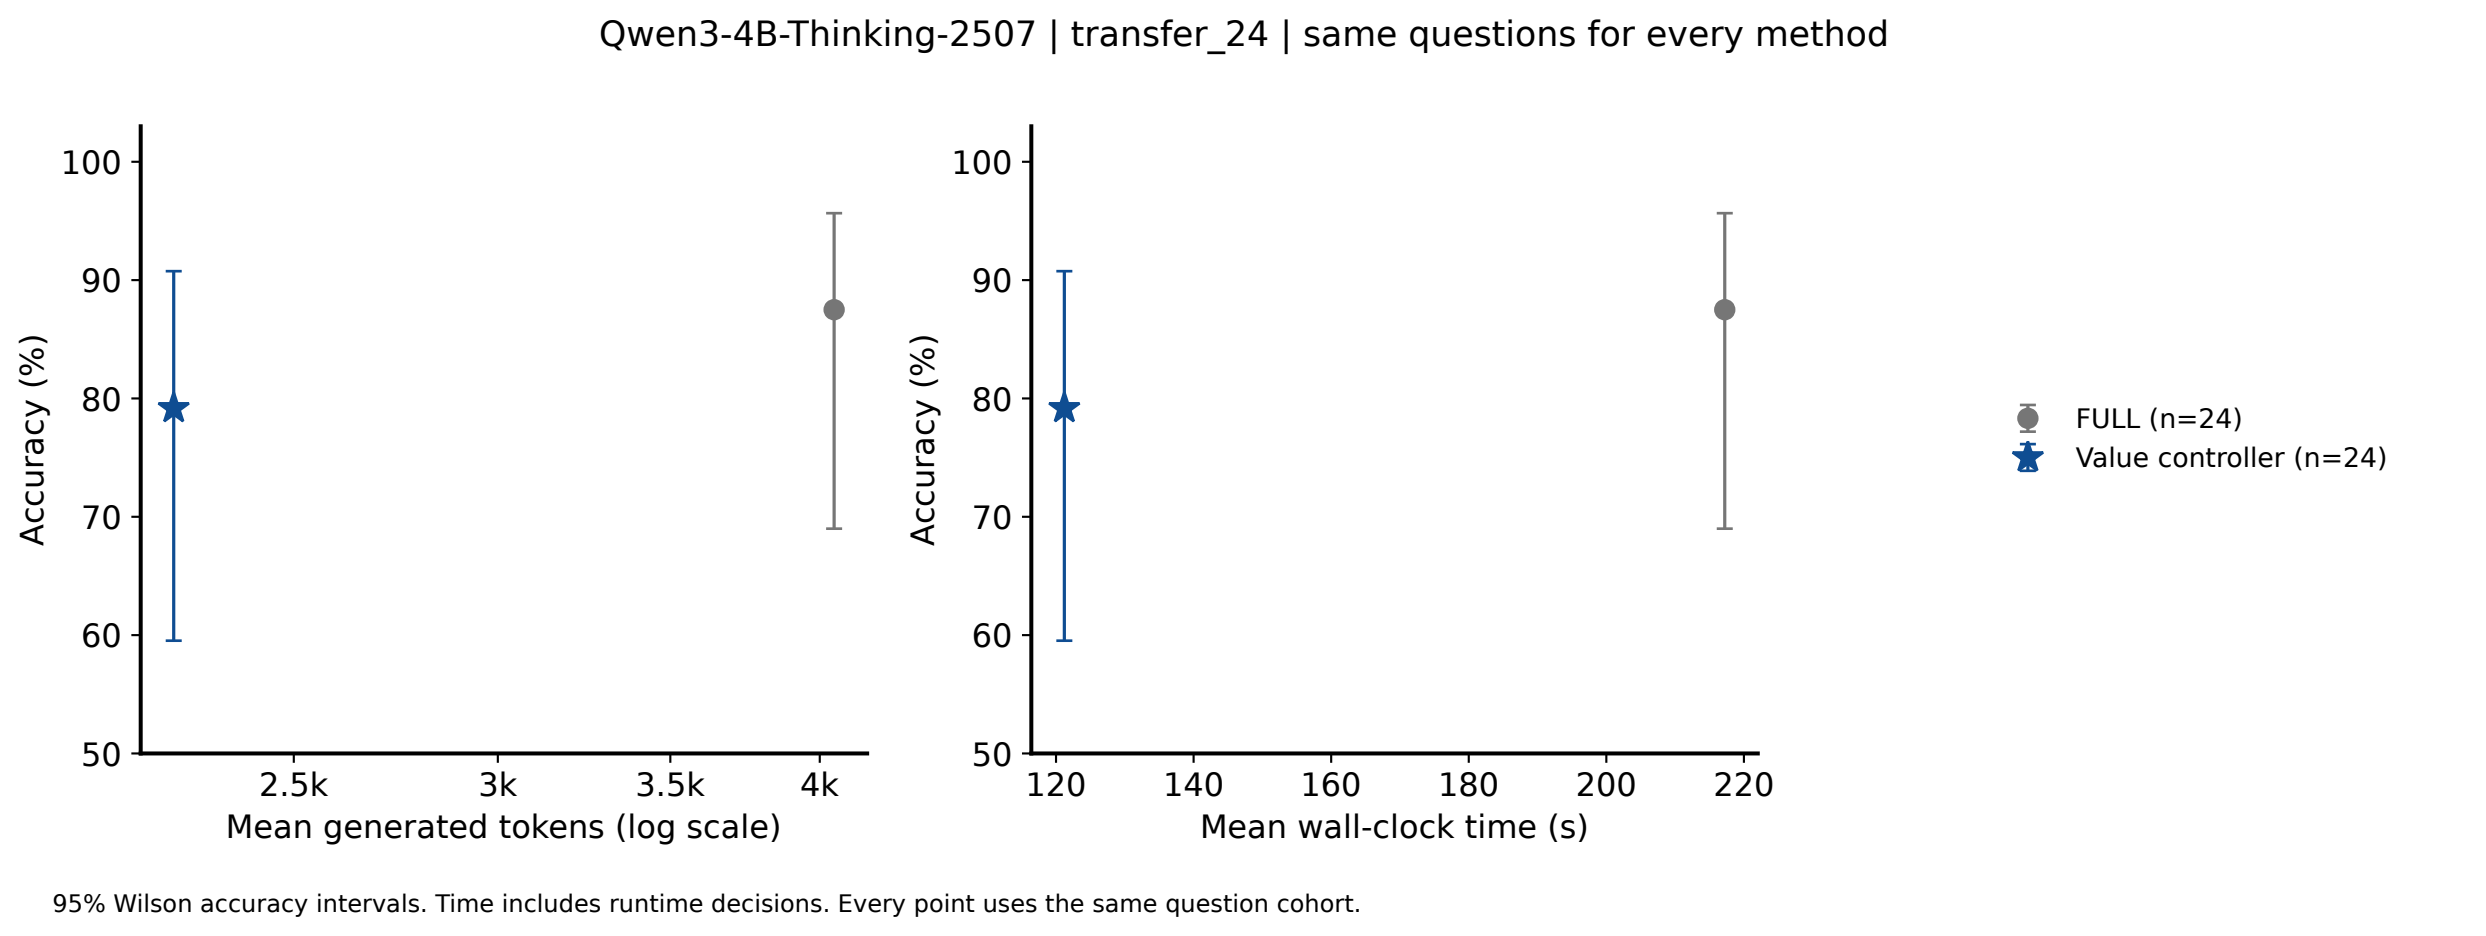
Qwen3-4B-Thinking-2507 | transfer_24 | same questions for every method
Accuracy (%)
100
90
80
70
60
50
2.5k
3k
3.5k
4k
Mean generated tokens (log scale)
Accuracy (%)
100
90
80
70
60
50
120
140
160
180
200
220
Mean wall-clock time (s)
FULL (n=24)
Value controller (n=24)
95% Wilson accuracy intervals. Time includes runtime decisions. Every point uses the same question cohort.