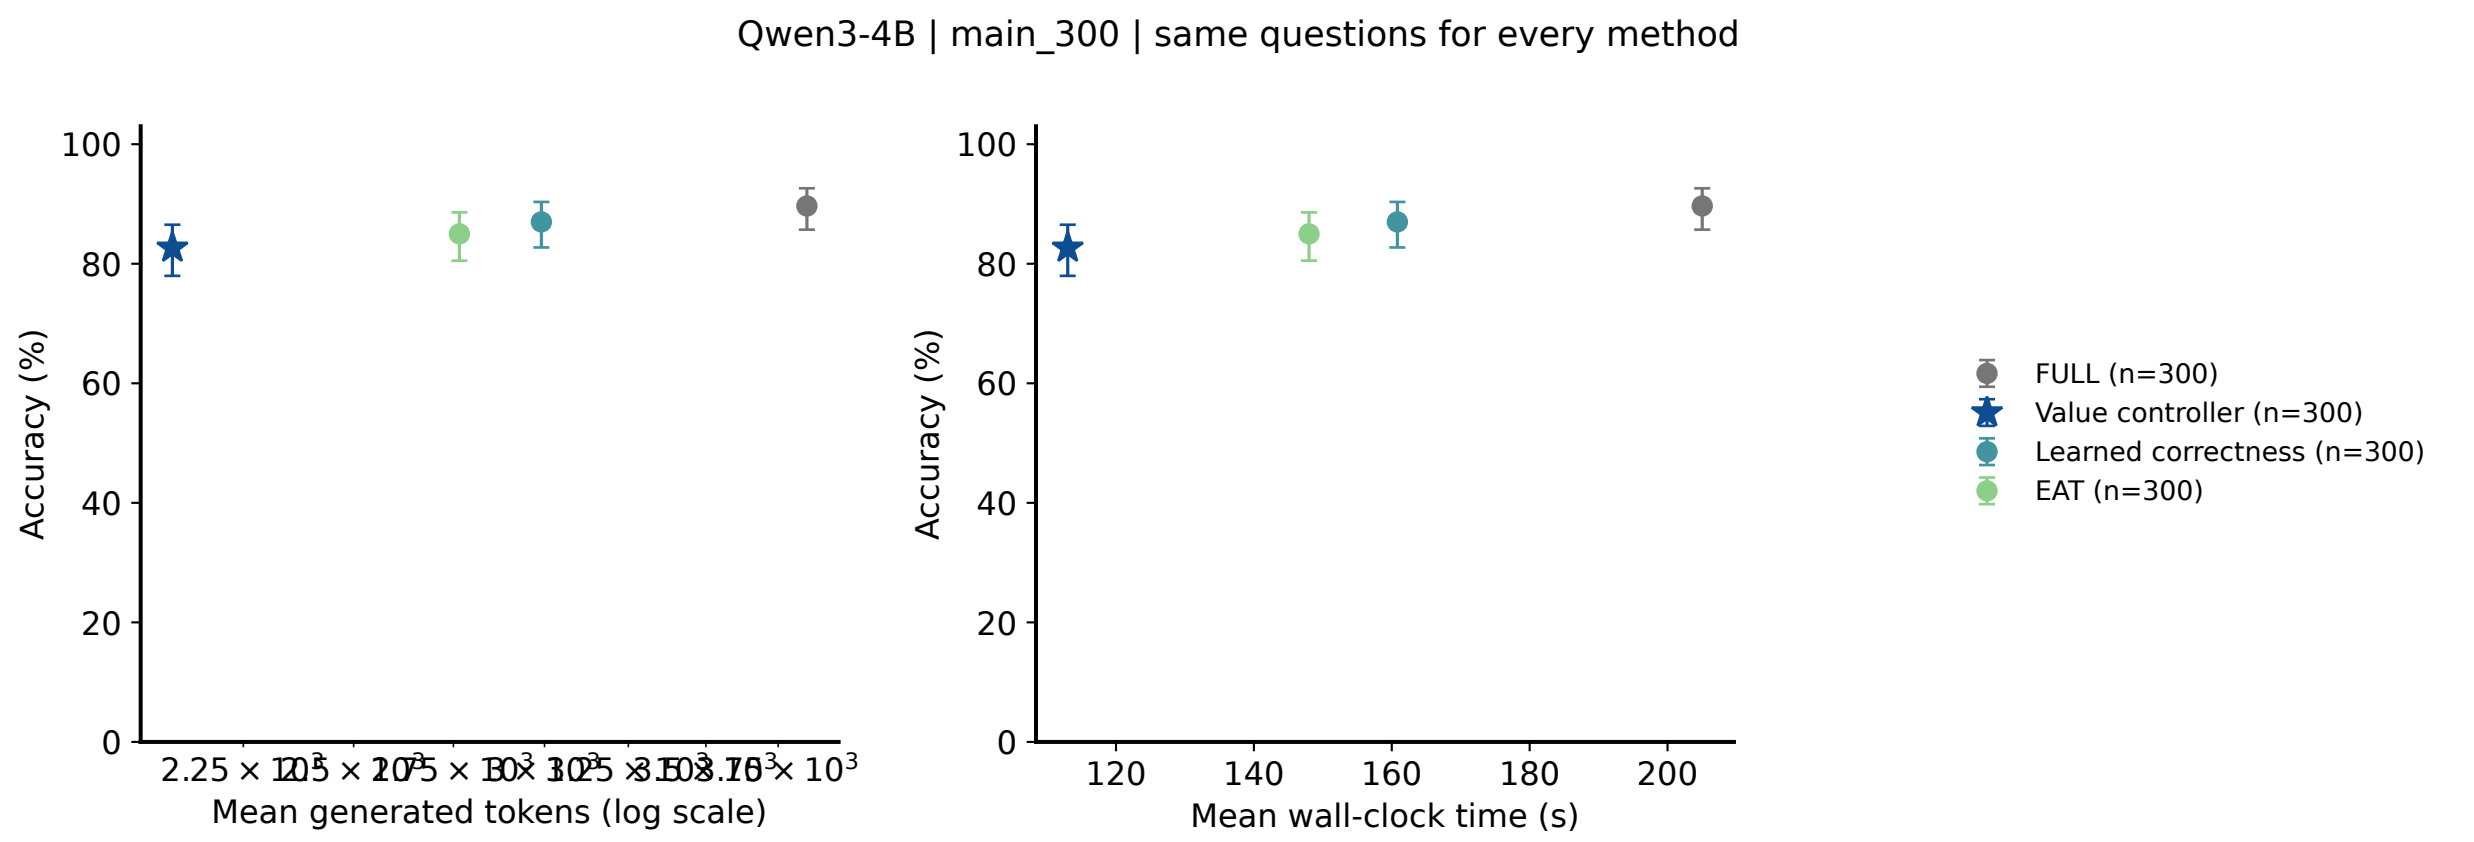
Qwen3-4B | main_300 | same questions for every method
Accuracy (%)
Mean generated tokens (log scale)
Mean wall-clock time (s)
FULL (n=300)
Value controller (n=300)
Learned correctness (n=300)
EAT (n=300)
95% Wilson accuracy intervals. Time includes runtime decisions. Every point uses the same question cohort.