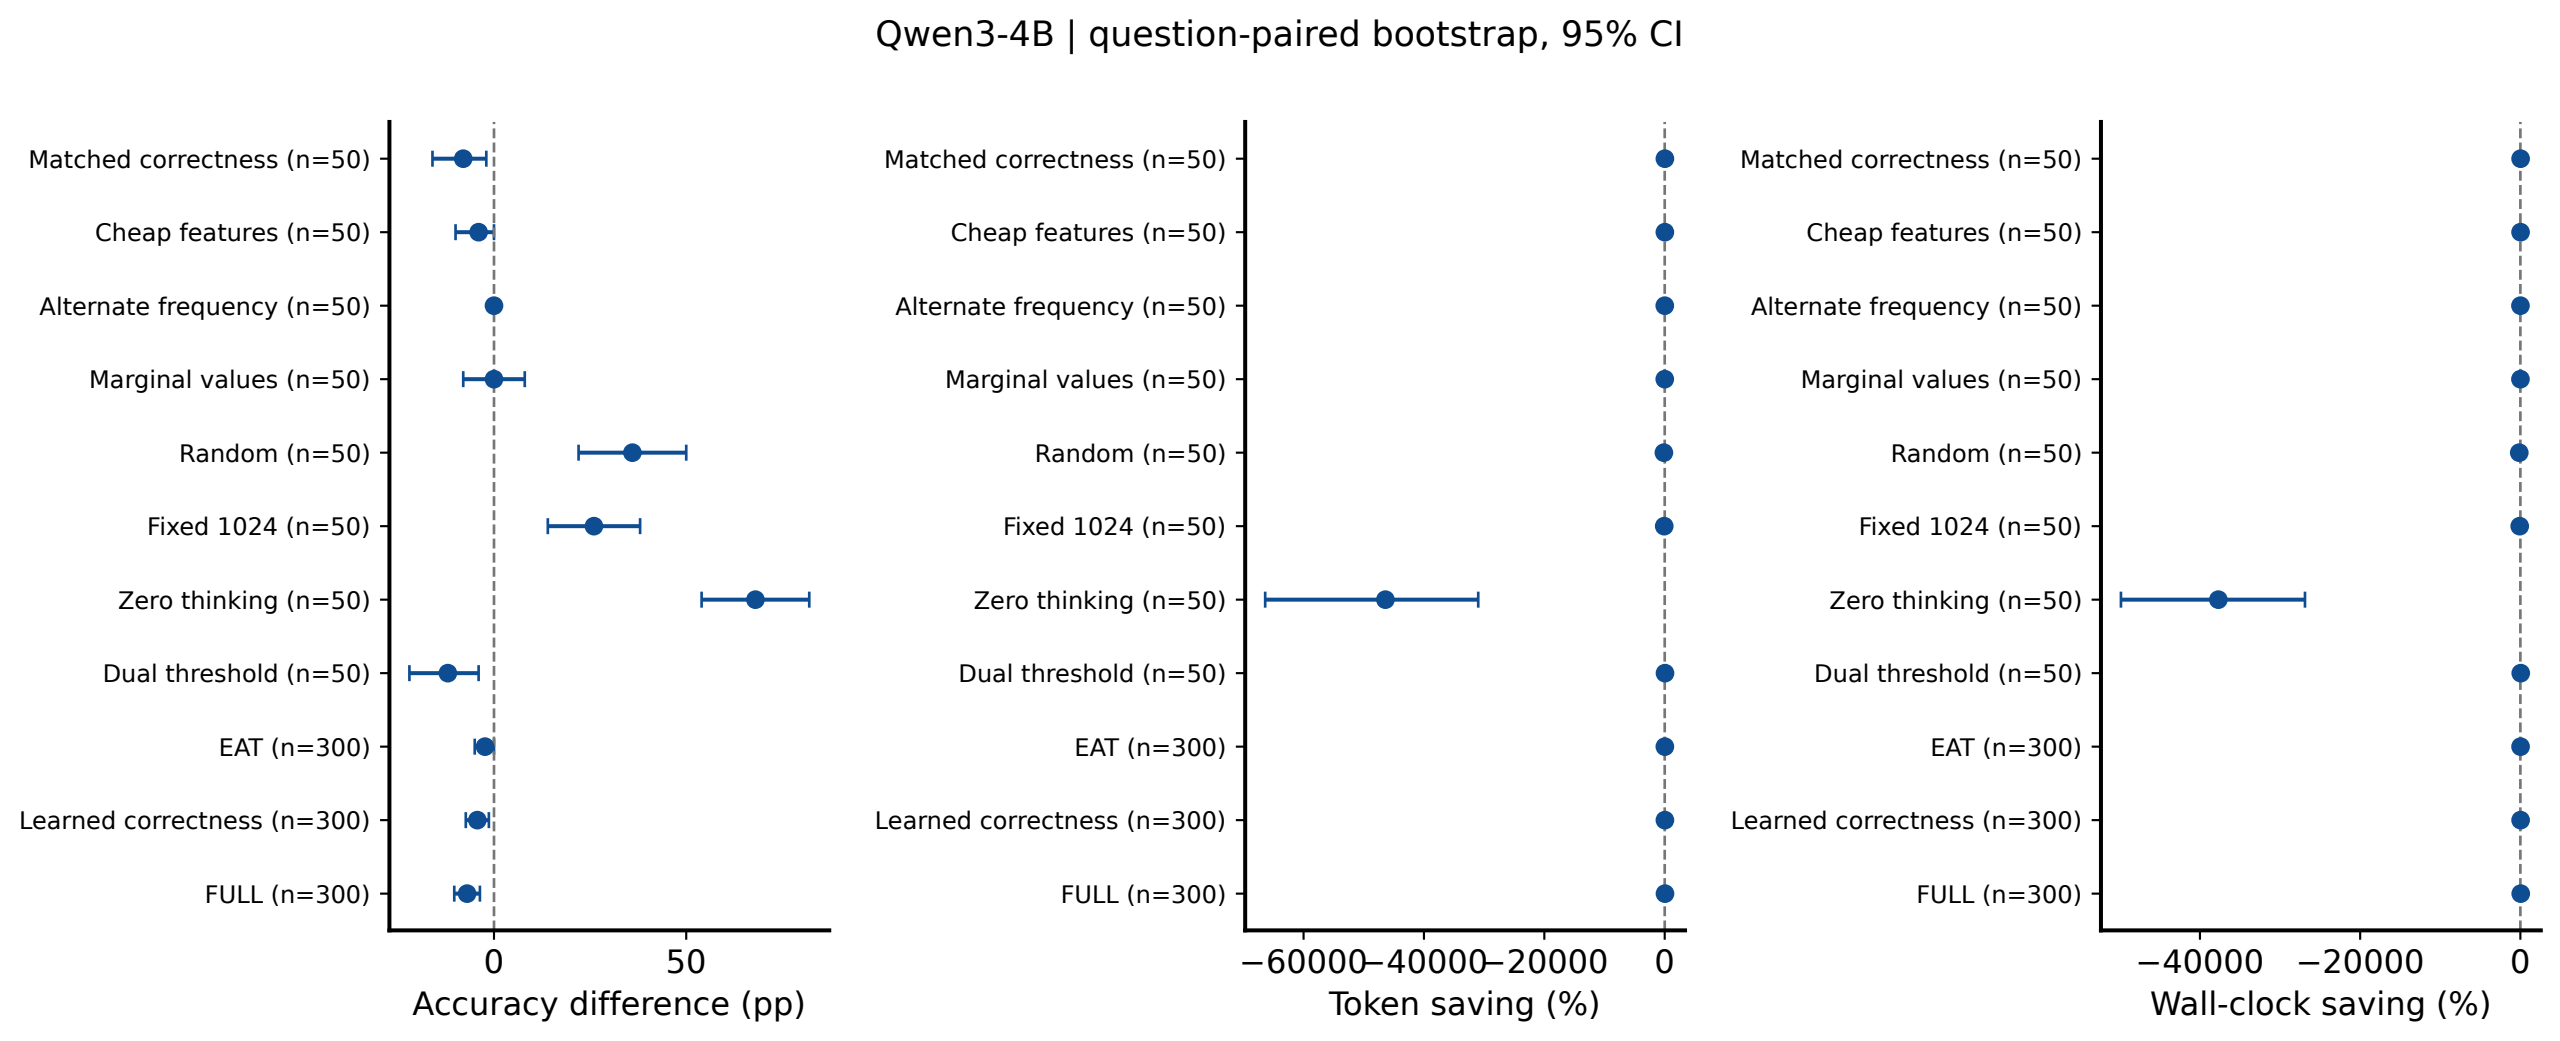
Qwen3-4B | question-paired bootstrap, 95% CI
Matched correctness (n=50)
Cheap features (n=50)
Alternate frequency (n=50)
Marginal values (n=50)
Random (n=50)
Fixed 1024 (n=50)
Zero thinking (n=50)
Dual threshold (n=50)
EAT (n=300)
Learned correctness (n=300)
FULL (n=300)
0
50
Accuracy difference (pp)
−60000
−40000
−20000
0
Token saving (%)
−40000
−20000
0
Wall-clock saving (%)